

Appendix M: Mathematical Formalisation of the Phase-Coherent Transformer

Companion document to Hioki,
Complex-Valued Phase-Coherent Transformer

Appendix M: Mathematical Formalisation of the Phase-Coherent Transformer

Companion document to: Hioki, *Complex-Valued Phase-Coherent Transformer*.

Status: Full proofs and explicit conditions. This is the rigorous version of the appendix; the paper’s main body contains a shorter intuitive sketch.

1 Mathematical Formalisation

This appendix contains the full mathematical formalisation of the two-level phase-coherence property (L1, L2) and the four-condition framework that grounds the closeness-to-PCT axis used in body §5.2. Body §5.2 provides a short overview; this appendix gives the formal definitions, theorem statements, and proofs.

1.1 Machine-checked proofs (Lean) – recommended

A Lean formalisation of the definitions and theorems in this appendix is available at:

<https://github.com/leohio/phase-coherent-transformer-r-d/tree/main/lean>

We recommend that readers who want a quick, mechanical sanity-check on the definitions and the sufficiency theorems consult the Lean development first. The Lean files mirror the structure of the appendix one-to-one:

- `PaperV4/Basic.lean` – Setting and notation (M.0): `TokenSeq`, the global rotation $R(\varphi)$ and per-token shift $P(\varepsilon)$, and the elementary identities used by Theorem 1 and Lemma A.
- `PaperV4/L1.lean` – Definition 1, Theorem 1 ($C1 + C4 \Rightarrow L1$), Theorem 1’, and Corollary 2. Machine-checked in full (no `sorry`).
- `PaperV4/LemmaA.lean` – Lemma A (M.7): the exact factorisation $P(\varepsilon) = R(\bar{\varphi}) \circ P(\delta)$ and the global-mode pass-through across an L1 stack. Machine-checked in full (no `sorry`); the quantitative $\|\tilde{Y}_L - Y_L\|$ bound is machine-checked as `lemmaA_bound` in `L2.lean` – exact, with no $O(\bar{\varphi}^2)$ remainder (via the global estimate $|e^{i\bar{\varphi}} - 1| \leq |\bar{\varphi}|$).
- `PaperV4/LemmaC.lean` – Lemma C (M.9, Doeblin contraction). Machine-checked in full (no `sorry`, 2026-08-05): the rank-1 Doeblin component annihilates the zero-mean subspace and the residual matrix has non-negative entries with row sums $1 - \mu_D$, so the finite-dimensional $\ell^\infty \rightarrow \ell^\infty$ bound follows by a direct big-operator argument – no coupling formalisation required.

- `PaperV4/L2.lean` – Definition 3, Definition 4, the quantitative Lemma A bound (`lemmaA_bound`), the M.6 step-4 geometric recurrence (`cascade_geometric_bound`), **Theorem5Hypotheses** (a `Prop` structure whose fields are exclusively standard mathematical conditions: per-layer L1.a, per-layer non-expansiveness, per-layer uniform zero-mean phase-Lipschitz stability, uniform output bound), and **theorem5** (Theorem 5 of M.10). **theorem5** is **proven without sorry**, with the depth-uniform constants *derived by induction* – no hypothesis is conclusion-shaped; `#print axioms PaperV4.theorem5` shows only the three standard Lean / Mathlib axioms [`propext`, `Classical.choice`, `Quot.sound`]. The hypothesis set is proven satisfiable (`theorem5Hypotheses_satisfiable`). Lemmas B (M.8) and D (M.10) are not re-proven in Lean against a concrete architecture – their single-layer conclusions are exactly the named standard hypotheses `zero_mean_phase_stable` / `substrate_nonexpansive`.

Status summary (verified 2026-08-05). Build is clean; **the project contains no sorry and no bespoke premises** (the earlier conclusion-shaped `Theorem5Premises.cascade_decomposition` field has been eliminated). Every hypothesis of every theorem is a community-standard condition (linearity, row-stochasticity, Doeblin minorisation, Lipschitz / non-expansiveness, uniform phase-Lipschitz stability, bounded range). Remaining Lean-side work – deriving the per-layer hypotheses for the concrete trained architecture – reduces to (i) Lemmas B and D (paper §M.5 classifies as “rigorous”, to be Lean’d), and (ii) the M.11 closure ($\Lambda_S \cdot \sup_l \|J_l|_{V_0}\| < 1$ preserved across layers + (S3) under training); this changes what is *instantiated*, not what is *assumed*.

The Lean development is intended as a low-friction verification path: a reader can clone the repository and run `lake build` to confirm that Theorem 1 and Theorem 5 type-check. Where the Lean and the prose diverge, the Lean version is the authoritative source for the assertion’s content.

Operating-range form of the conditions. The L2-normalisation on Q, K restricts the cosine score to the operating range $s \in [-d, d]$, so the gate $\alpha = f$ only ever sees inputs from `[-sqrtd - |b|, sqrtd + |b|]` – call this the **operating range** `D_op`. We adopt the convention that C1, C2, C3, C4 are stated in their operating-range form. This convention is *load-bearing*:

- A cell that is C2-violating on $\mathbb{R}$ but bounded on `D_op` **satisfies operating-range C2** and is empirically close to PCT. The substrate **bypasses** the strict-on- $\mathbb{R}$ violation.
- A cell whose violation persists on `D_op` **fails operating-range C2** and exhibits cascade degradation.

Theorem 5 below requires `C1 + C2 + C3 + C4` in the operating-range form, and the empirical results in §5.4 confirm the necessity of all four conditions in this form. The proof’s M constant is taken as the operating-range bound throughout.

Three structural grades of deviation. A cell can deviate from a condition `C_k` in three structurally distinct ways:

1. **Partial in the gate domain:** `C_k` holds on a subset of `D_op` and is violated on the complement.
2. **Bypassed by the substrate:** `C_k` violated on $\mathbb{R}$ but operating-range form holds, courtesy of L2-normalisation bounding `D_op`.
3. **Strict in the operating range:** `C_k` violated on most or all of `D_op`, not bypassed.

Theorem 5 holds for cells whose deviations are *partial* or *bypassed*. Cells with *strict-in-operating-range* deviations fall outside Theorem 5’s hypotheses and exhibit empirical collapse, as confirmed by the §5.4 isolation.

1.2 M.0 Setting and notation

- Token sequence $X \in \mathbb{C}^{N \times d}$, with $x_i \neq 0$.
- Global phase shift $R(\varphi) : X \mapsto (e^{i\varphi}x_1, \dots, e^{i\varphi}x_N)$ for $\varphi \in \mathbb{R}$.
- Per-token phase shift $P(\varepsilon) : X \mapsto (e^{i\varepsilon_i}x_i)_i$ for $\varepsilon \in \mathbb{R}^N$.
- A complex attention layer $A_\theta : \mathbb{C}^{N \times d} \rightarrow \mathbb{C}^{N \times d}$ is parameterised by complex-linear maps W_q, W_k, W_v, W_o and a gate function $f : \mathbb{R} \rightarrow \mathbb{R}$.
- Standard form: writing $q_i = W_q x_i, k_j = W_k x_j, v_j = W_v x_j, \bar{q}_i = q_i / \|q_i\|_2, \bar{k}_j = k_j / \|k_j\|_2$, the layer is

$$\begin{aligned} s_{ij}(X) &= \text{Re}\langle \bar{q}_i, \bar{k}_j \rangle \quad (\text{cosine score}, \in [-1, 1]) \\ \alpha_{ij}(X) &= f(s_{ij}(X) + b) \quad (\text{gate}, \in \mathbb{R}) \\ A_\theta(X)_i &= W_o \cdot \left(\sum_j \alpha_{ij}(X) \cdot v_j \right) \end{aligned}$$

1.3 M.1 Per-layer phase-coherence (L1) – definition

Definition 1. An attention layer $A : \mathbb{C}^{N \times d} \rightarrow \mathbb{C}^{N \times d}$ is **per-layer phase-coherent** if it satisfies both:

(L1.a) Global phase equivariance. For every $\varphi \in \mathbb{R}$ and every X , $A(R(\varphi)X) = R(\varphi) \cdot A(X)$.

(L1.b) Element-independent gating. There exist a real-valued function $f : \mathbb{R} \rightarrow \mathbb{R}$, a complex-linear value path $V : \mathbb{C}^d \rightarrow \mathbb{C}^d$, and a binary score function $s : \mathbb{C}^d \times \mathbb{C}^d \rightarrow \mathbb{R}$ invariant under joint global phase rotation ($s(e^{i\varphi}a, e^{i\varphi}b) = s(a, b)$), such that the layer can be written $A(X)_i = \sum_j f(s(x_i, x_j)) \cdot V(x_j)$.

(L1.b) is the formal version of "no row-norm coupling": each gate value α_{ij} is determined by the pair (x_i, x_j) alone, with no dependence on tokens x_k for $k \notin i, j$.

1.4 M.2 L1 sufficiency – theorem

Theorem 1. Suppose an attention layer A_θ is constructed with complex-linear maps W_q, W_k, W_v, W_o and a real-valued gate $\alpha_{ij} = f(s_{ij})$ where $s_{ij} = \text{Re}\langle \bar{q}_i, \bar{k}_j \rangle$. Then A_θ satisfies Definition 1.

Proof. For (L1.a): under $X \mapsto R(\varphi)X$, complex-linearity gives $q_i \mapsto e^{i\varphi}q_i$, hence $\bar{q}_i \mapsto e^{i\varphi}\bar{q}_i$; similarly $\bar{k}_j \mapsto e^{i\varphi}\bar{k}_j$. Then $\langle e^{i\varphi}\bar{q}_i, e^{i\varphi}\bar{k}_j \rangle = (e^{i\varphi})^* \cdot e^{i\varphi} \cdot \langle \bar{q}_i, \bar{k}_j \rangle = \langle \bar{q}_i, \bar{k}_j \rangle$, so s_{ij} , hence α_{ij} , is invariant. Meanwhile $v_j \mapsto e^{i\varphi}v_j$, so $\sum_j \alpha_{ij}v_j \mapsto e^{i\varphi}\sum_j \alpha_{ij}v_j$, and W_o is complex-linear, so $A_\theta(R(\varphi)X) = R(\varphi)A_\theta(X)$. ✓

For (L1.b): take f as the gate function, $V = W_o \cdot W_v$ as the value path, and $s(a, b) = \text{Re}\langle \bar{a}, \bar{b} \rangle$. Then $f(s(x_i, x_j))$ involves only (x_i, x_j) , and $s(e^{i\varphi}a, e^{i\varphi}b) = s(a, b)$ by the same conjugate-cancellation as above. ✓ ■

Corollary 2.

- *PCT*: gate is $f(s) = \sigma$, real-valued (C1); element-wise dependent on s_{ij} only (C4). By Theorem 1, PCT is L1-coherent.

- **complex_screen**: gate is $f(s) = T(r^2(s) \cdot \max(s - t, 0)^2)$ where $T = \text{TanhNorm}$ is per-token, r is the magnitude prefactor depending on $(\mathbf{x}_i, \mathbf{x}_j)$ only, and t is a learnable scalar. The gate is real-valued (C1) and per-pair (C4). By Theorem 1, **complex_screen** is L1-coherent.
- **complex_softmax**: gate is $\alpha_{ij} = \exp(s_{ij}) / \sum_k \exp(s_{ik})$, which depends on $s_{ik} : k \in [N]$. (L1.b) fails: there is no representation $\alpha_{ij} = f(s_{ij})$ because the denominator couples tokens.

Theorem 1’. If a gate of the form $\alpha_{ij} = f$ for some f is L1-coherent in the sense of Definition 1, then f must be expressible as a function of $\mathbf{s}_{ij}$ alone for each pair.

Proof sketch. L1.b requires $\alpha_{ij} = f(s_{ij})$; this is precisely the negation of C4-violating row coupling. ■

This closes the L1 layer. Per-layer phase-coherence is captured exactly by C1 (real) + C4. C2 and C3 do not enter L1.

1.5 M.3 All-layer phase-coherence (L2) – definition

For L2 we capture that *stacking* L per-layer-coherent layers preserves a phase invariant that does not degrade with L . We formalise this as **cascade phase stability**: the composed map should be Lipschitz-stable under per-token phase perturbations of the input, with a Lipschitz constant independent of L .

Definition 3. A composition $A_L \circ A_{L-1} \circ \dots \circ A_1$ of per-layer-coherent attention layers is **cascade phase stable** if there exist constants $C_0, C_1 > 0$, *independent of L* , such that for every input X and every per-token phase perturbation $\varepsilon \in \mathbb{R}^N$ with $\|\varepsilon\|_\infty \leq \delta$,

$$\|(A_L \circ \dots \circ A_1)(P(\varepsilon)X) - (A_L \circ \dots \circ A_1)(X)\|_2 \leq C_0 \cdot \delta + C_1 \cdot \delta^2.$$

The key requirement is *L-independence* of C_0, C_1 . A trivial bound of the form $(1 + \gamma)^L \cdot \delta$ is *not* cascade phase stability – it allows phase noise to grow exponentially with depth.

Definition 4. A stack of attention layers is **all-layer phase-coherent** if it (i) is per-layer phase-coherent at every layer and (ii) is cascade phase stable.

1.6 M.4 L2 sufficiency – refined conjecture

Conjecture 5 (refined: C1+C3+C4 + substrate $\Rightarrow$ L2). Suppose every layer in the stack satisfies:

- (C1) Real-valued gate $\alpha_{ij} \in \mathbb{R}$.
- (C3) Lipschitz gate function $f : \mathbb{R} \rightarrow \mathbb{R}$ with $|f'(s)| \leq K_f$ for all $s \in \mathbb{R}$, and $f' > 0$ on a positive-measure subset of $\mathbb{R}$.
- (C4) Element-independent gate $\alpha_{ij} = f(s_{ij})$.

Suppose further the architectural baseline:

- (A1) **L2-normalize on Q, K** so cosine score $s_{ij} \in [-d, d]$ regardless of input.
- (A2) **Continuous gate f on operating range** so $\|\alpha\|_\infty \leq M := \sup_{s \in [-d-|b|, d+|b|]} |f|$ is finite.

Suppose further the substrate is bounded:

- (S1) Value path operator norm $\|W_v\|_{op} \cdot \|W_o\|_{op} \leq V_m ax$.
- (S2) Layer-wise residual + normalisation makes the per-layer Lipschitz constant of the composed sub-map at most $\Lambda < \infty$.

Then there exist constants $C_0, C_1 > 0$ depending only on $(M, K_f, V_m ax, \Lambda, N, d)$ such that the stack is cascade phase stable with bound $C_0\delta + C_1\delta^2$, *independent of L* .

Status: drafted in M.6–M.10. Two residual technical pieces flagged in M.11.

Note on the C2 $\rightarrow$ (A2) reformulation: The original C2 is **redundant** in the L2-normalized PCT architecture. (A1) provides bounded score domain, and any continuous gate (A2) automatically gives bounded $\|\alpha\|_\infty$ on the operating range. The §5.4 isolation experiment (**complex_softplus** violates C2 strict-on- $\mathbb{R}$ but satisfies (A2), and achieves acc=1.000 on Copy d=1000 N=3) confirms that C2 strict-on- $\mathbb{R}$ is not the operative condition; what matters is operating-range boundedness which is automatic.

1.7 M.5 Status: rigorous vs conjectured

Rigorous:

- Definition 1 (L1) and Definition 3 – full statements above.
- Theorem 1 (C1 + C4 $\Rightarrow$ L1) – full elementary proof above.
- Corollary 2 – direct from Theorem 1 / Theorem 1’.
- Lemma A – full proof.
- Lemma B – full algebraic derivation; second-order remainder bounded but not tightened.
- Lemma C – full proof, citing Levin–Peres–Wilmer 2017 Theorem 4.9 for the standard Doeblin coupling argument.

Drafted with explicit conditions, full proof modulo two residual pieces (M.11):

- Theorem 5 – proven as the composition of Lemmas A–D plus Corollary C.1.
- Corollary C.1 – proof outlined; the cross-token term inherits Doeblin contraction directly (rigorous), the self-token / residual diagonal term is bounded but not contractive (rigorous), and the closure condition $K_R < \mu_D$ requires the bounded-input regime to be preserved across layers.
- Lemma D – standard transformer-stability argument citing Wang–Sun 2023 (DeepNet); rigorous under spectral-normalisation or weight-decay assumptions, otherwise empirical.

Two residual technical pieces (M.11) – both tractable, neither is a deep open problem:

1. Empirical verification of (S3) – attention diffuseness preserved across training, with μ_D quantified per layer for trained PCT checkpoints.
2. Fixed-point argument that the bounded-input regime where $K_R < \mu_D$ is preserved across the cascade.

C2 status: The original C2 is redundant; replaced by (A2) "continuous gate on bounded operating range", automatic in PCT architecture.

1.8 M.6 Strategy of the proof of Conjecture 5

The strategy decomposes the per-token phase perturbation $\varepsilon \in \mathbb{R}^N$ into two orthogonal modes and bounds each separately:

1. **Global-mode decomposition.** Write $\varepsilon = \varphi \cdot 1 + \delta$ where $\varphi = (1/N)\sum_i \varepsilon_i$ and $\delta \perp 1$ (so $\sum_i \delta_i = 0$). The global mode $\varphi \cdot 1$ passes through every layer *exactly* by L1.a. Hence the cascade error from the global mode is $O(\delta) \cdot \|\text{output}\|$, which is L-independent provided the stack output norm is bounded.
2. **Linearisation on the zero-mean subspace.** A single layer's effect on a zero-mean phase perturbation δ is, to first order in $\|\delta\|_\infty$, given by a Jacobian matrix $J_l \in \mathbb{R}^{N \times N}$ whose action depends on the row-stochasticised gate $P_{ij} = \alpha_i j / \sum_k \alpha_i k$ and the cosine-score's imaginary part $\eta_i j = -\text{Im}\langle \bar{q}_i, \bar{k}_j \rangle$.
3. **Doebelin contraction on zero-mean subspace (Lemma C, proven assuming (S3)).** Under attention diffuseness (S3) – every gate value is bounded below by a strictly positive multiple of a stationary distribution π – the row-stochastic matrix $P_1 = (P_{ij})^{(1)}$ satisfies a Doebelin condition with explicit constant $\mu_D > 0$. This implies the operator norm of J_1 restricted to the zero-mean subspace $\delta : \sum_i \delta_i = 0$ is bounded above by $\Lambda := 1 - \mu_D < 1$.
4. **Cascade summation.** The depth-L cascade Lipschitz on the zero-mean subspace is bounded geometrically by $\sum_{l=0}^{L-1} \Lambda^l \leq 1/(1 - \Lambda) = 1/\mu_D$, *L-independent*.
5. **Substrate non-expansion.** Residual + RMSNorm + FFN composition is non-expansive on bounded-norm inputs.

The proof is complete modulo:

- (i) **Verification of (S3)** for trained PCT layers: at initialisation (S3) holds with $\mu = e^{-1}/N$ and $\pi = 1/N \cdot 1$. Whether (S3) is preserved across training requires either an empirical check or a stability-of-training argument.
- (ii) **Tightening the second-order term:** the $O(\delta^2)$ bound is loose; a careful Taylor-2 expansion is straightforward but technical.

1.9 M.7 Lemma A – Global-mode decomposition

Lemma A. Let $\varepsilon \in \mathbb{R}^N$ and write $\varepsilon = \varphi \cdot 1 + \delta$ with $\varphi = (1/N)\sum_i \varepsilon_i$ and $\sum_i \delta_i = 0$. Let $Y_L := (A_L \circ \dots \circ A_1)(X)$ and $\tilde{Y}_L := (A_L \circ \dots \circ A_1)(P(\varepsilon)X)$. Suppose every A_l is per-layer phase-coherent. Then

$$\tilde{Y}_L = R(\varphi) \cdot (A_L \circ \dots \circ A_1)(P(\delta)X),$$

and consequently

$$\|\tilde{Y}_L - Y_L\|_2 \leq \|(A_L \circ \dots \circ A_1)(P(\delta)X) - Y_L\|_2 + |\varphi| \cdot \|Y_L\|_2 + O(\varphi^2).$$

Proof. The exact factorisation $P(\varepsilon) = R(\varphi) \circ P(\delta)$ holds because $e^{i\varepsilon_i} = e^{i\varphi} \cdot e^{i\delta_i}$. By L1.a applied iteratively, $A_l \circ R(\varphi) = R(\varphi) \circ A_l$ for every l , so

$$\tilde{Y}_L = (A_L \circ \dots \circ A_1)(R(\varphi)P(\delta)X) = R(\varphi) \cdot (A_L \circ \dots \circ A_1)(P(\delta)X).$$

With $Z := (A_L \circ \dots \circ A_1)(P(\delta)X)$ and $Y_L = (A_L \circ \dots \circ A_1)(X)$:

$$\|\tilde{Y}_L - Y_L\| = \|R(\varphi)Z - Y_L\| \leq \|R(\varphi)(Z - Y_L)\| + \|(R(\varphi) - I)Y_L\| = \|Z - Y_L\| + \|(R(\varphi) - I)Y_L\|$$

(using that $R(\varphi)$ is unitary). For $R(\varphi)y_i = e^{i\varphi}y_i$, we have $|(e^{i\varphi} - 1)y_i| \leq |\varphi| \cdot |y_i| + O(\varphi^2)$. Summing, $\|(R(\varphi) - I)Y_L\|_2 \leq |\varphi| \cdot \|Y_L\|_2 + O(\varphi^2)$. ■

The first term is the cascade Lipschitz under **zero-mean** phase perturbation. The second term is L-independent because $\|Y_L\|$ is bounded under (S1)+(S2).

1.10 M.8 Lemma B – Linearised Jacobian on the zero-mean subspace

Lemma B. *Fix a layer A_l with parameters θ and gate f . For every input X and every zero-mean $\delta \in \mathbb{R}^N$ with $\|\delta\|_\infty \leq \xi$ (small), the layer output satisfies*

$$A_l(P(\delta)X) = A_l(X) + i \cdot J_l(X) \cdot \delta + O(\xi^2)$$

where $J_l(X)$ decomposes as

$$J_l(X) \cdot \delta = T_1 \cdot \delta - T_2 \cdot \delta + T_3 \cdot \delta + \rho \cdot \text{diag}(x) \cdot \delta$$

with

- $[T_1\delta]_i := W_o \cdot \Sigma_j \beta_{ij}^{(l)} \delta_j v_j$,
- $[T_2\delta]_i := \delta_i \cdot W_o \cdot \Sigma_j \beta_{ij}^{(l)} v_j$,
- $[T_3\delta]_i := W_o \cdot \Sigma_j \alpha_{ij}^{(l)} \delta_j v_j$,
- $[\rho \cdot \text{diag}(x)\delta]_i := \rho \cdot \delta_i \cdot x_i$,

where $\beta_{ij} = f' \cdot \eta_{ij}$, $\eta_{ij} = -\text{Im}\langle \bar{q}_i, \bar{k}_j \rangle$, $|\beta_{ij}| \leq K_f$, $|\alpha_{ij}| \leq M$.

Proof sketch. Write $\bar{q}'_i = e^{i\delta_i} \bar{q}_i + O(\xi^2)$, similarly $\bar{k}'_j, v'_j$. Then

$$s'_{ij} = \text{Re}\langle \bar{q}'_i, \bar{k}'_j \rangle = \text{Re}(e^{i(\delta_j - \delta_i)} \langle \bar{q}_i, \bar{k}_j \rangle).$$

Linearising and using $\langle \bar{q}_i, \bar{k}_j \rangle = s_{ij} + i \cdot (-\eta_{ij})$:

$$s'_{ij} - s_{ij} = \eta_{ij}(\delta_j - \delta_i) + O(\xi^2).$$

Hence $\alpha'_{ij} - \alpha_{ij} = \beta_{ij}(\delta_j - \delta_i) + O(\xi^2)$. The value: $v'_j = e^{i\delta_j} v_j = v_j + i\delta_j v_j + O(\xi^2)$. Combining:

$$A_l(P(\delta)X)_i = W_o \cdot \Sigma_j [\alpha_{ij} + \beta_{ij}(\delta_j - \delta_i)] [v_j + i\delta_j v_j] + \rho(1 + i\delta_i)x_i + O(\xi^2),$$

from which $T_1, T_2, T_3, \rho \cdot \text{diag}(x)$ are read off after splitting $\Sigma_j \beta_{ij}(\delta_j - \delta_i) = \Sigma_j \beta_{ij}\delta_j - \delta_i \Sigma_j \beta_{ij}$. ■

1.11 M.9 Lemma C – Doeblin contraction on the zero-mean subspace

Lemma C. *Suppose the gate satisfies (C1)–(C4) and (S3) – there exist $\mu > 0$ and a probability vector π such that $\alpha_{ij} \geq \mu \cdot \pi_j$ for all i, j and all inputs in the bounded set. Define the row-stochastic matrix P by $P_{ij} := \alpha_{ij} / \Sigma_k \alpha_{ik}$. Then P satisfies the Doeblin condition*

$$P_{ij} \geq \mu_D \cdot \pi_j \text{ with } \mu_D := \mu/M,$$

and the operator norm of P restricted to the zero-mean subspace $V_0 := \{u \in \mathbb{R}^N : \Sigma_i u_i \cdot w_i = 0\}$ is bounded by

$$\|P|_{V_0}\|_\infty \rightarrow \infty \leq 1 - \mu_D.$$

Proof. The Doeblin condition follows from $P_{ij} \geq \mu\pi_j/M$. The contraction on the zero-mean

subspace is the standard *coupling lemma*: given $P_{ij} \geq \mu_D \pi_j$, write $P = \mu_D \cdot 1\pi^T + (1 - \mu_D) \cdot Q$ for some row-stochastic Q ; the rank-1 component annihilates $\{u : \langle \pi, u \rangle = 0\}$, and Q is non-expansive. ■

Corollary C.1. *Under (C1)–(C4), (S1)–(S3), the per-layer Jacobian J_l of Lemma B, restricted to the zero-mean subspace V_0 (uniform $\pi = 1/N \cdot 1$), has operator norm*

$$\|J_l|_{V_0}\|_{2 \rightarrow 2} \leq \Lambda_l \cdot W_{max}^2 \cdot R \cdot N$$

for some $\Lambda_l \in [0, 1)$.

The careful argument transferring P 's Doeblin contraction to J_l 's contraction on V_0 defines a phase-coordinate projection $\Pi : \mathbb{C}^{N \times d} \rightarrow \mathbb{R}^N$ and shows $\Pi J_l|_{V_0}$ is Doeblin-contractive. Cross-token terms T_1 , T_3 inherit Doeblin contraction directly; self-token terms T_2 and the residual diagonal are bounded but not contractive on their own – the contraction returns at the *next* attention call. Full detailed argument in M.11.

1.12 M.10 Lemma D – Substrate non-expansion + Theorem 5

Lemma D. *Suppose each layer A_l is followed by a residual + RMSNorm + FFN substrate S_l with the standard pre-norm transformer block structure. Under (S1) and bounded-input regime, S_l is non-expansive on bounded-norm phase perturbations: $\|S_l(\Delta u)\|_2 \leq \Lambda_S \cdot \|\Delta u\|_2$ for some $\Lambda_S \leq 1$.*

Proof. Standard. RMSNorm is 1-Lipschitz on bounded inputs; residual + FFN with bounded weights has standard transformer-stability bounds. ■

Theorem 5. *Suppose each layer A_l satisfies (C1)–(C4), is followed by a substrate S_l satisfying Lemma D, and (S1)–(S3) hold uniformly across layers. Then there exist constants C_0 , C_1 , **independent of L** , such that for all inputs X and all $\varepsilon \in \mathbb{R}^N$ with $\|\varepsilon\|_\infty \leq \delta$,*

$$\|(A_L \circ S_{L-1} \circ \dots \circ S_1 \circ A_1)(P(\varepsilon)X) - (A_L \circ \dots \circ A_1)(X)\|_2 \leq C_0 \cdot \delta + C_1 \cdot \delta^2.$$

Proof. By Lemma A, the global-mode contribution is $C_{global} \cdot \delta$ with C_{global} L -independent. For the zero-mean cascade, iterate Corollary C.1 + Lemma D:

$$\|\Delta u^{(l)}\| \leq \Lambda \cdot \|\Delta u^{(l-1)}\| + O(\delta^2)$$

where $\Lambda := \Lambda_S \cdot \sup_l \|J_l|_{V_0}\| \leq \Lambda_S \cdot (1 - \mu_D)$. If $\Lambda < 1$, the geometric series gives

$$\|\Delta u^{(L)}\| \leq \|\Delta u^{(0)}\|/(1 - \Lambda) + O(\delta^2)/(1 - \Lambda),$$

L -independent. Combining,

$$LHS \leq \delta + C_{quadratic} \cdot \delta^2 = C_0 \cdot \delta + C_1 \cdot \delta^2. \quad \blacksquare$$

For PCT, $\Lambda < 1$ is achievable. For `complex_relu`, f is C^0 not C^1 – the linearisation in Lemma B does not hold uniformly through the discontinuity, and the cascade summation does not close. **This is exactly the L2 failure.**

1.13 M.11 What remains rigorously open

Two residual technical pieces:

1. **Verification of (S3) under training dynamics.** We assumed $\alpha_i j \geq \mu \cdot \pi_j$ uniformly. At

initialisation, the $b = -\log N$ bias gives $\mu \geq e^{-1}/N$ cleanly. During training, gradient updates can sharpen attention; we conjecture (S3) holds across training. Empirical verification: scatter $\min_{i,j} \alpha_{ij}^{(l)}$ across training steps for trained PCT checkpoints.

2. **Step (2) of Corollary C.1:** transferring Doeblin contraction from P to J_1 |_{V_0}. The cross-token terms T_1, T_3 inherit Doeblin contraction directly; the self-token terms require the phase-coordinate projection argument detailed below.

Detailed argument for self-token term (2). Define $\Pi : \mathbb{C}^{N \times d} \rightarrow \mathbb{R}^N$ by $\Pi(u)_i := \text{Im}\langle u_i, \Delta u_i \rangle / \|u_i\|^2$. For a perturbation $\Delta u = J_l \cdot \delta$ arising from a zero-mean phase perturbation of input, $\Pi(\Delta u)$ is the induced output phase perturbation.

Apply Π to each term:

- $\Pi(T_1 \delta)_i = \text{Im}\langle u_i, [T_1 \delta]_i \rangle / \|u_i\|^2 = \text{Im}\langle u_i, W_o \Sigma_j \beta_{ij} \delta_j v_j \rangle / \|u_i\|^2$. Doeblin applies when $\beta_{ij} v_j \geq \mu'_D \cdot \pi_j \cdot v_j$ for some $\mu'_D > 0$ – valid under (S3).
- $\Pi(T_2 \delta)_i = \delta_i \cdot \text{Im}\langle u_i, \gamma_i \rangle / \|u_i\|^2$. Diagonal operator on δ ; bounded but not contractive. Contraction returns at the next attention call.
- $\Pi(T_3 \delta)_i$: same structure as T_1 with α_{ij} weights, Doeblin directly under (S3).
- $\Pi(\rho \cdot \text{diag}(x) \delta)_i = \rho \cdot \delta_i \cdot \text{Im}\langle u_i, x_i \rangle / \|u_i\|^2$: diagonal again, similarly bounded.

The **net effect** on $\Pi(\Delta u)$ is δ where $\mathcal{C}$ is the cross-token contraction ($\|C|_{V_0}\| \leq 1 - \mu_D$) and $\mathbf{R}$ is the diagonal residual ($\|R\|_\infty \leq K_R$). On V_0 : $\|v_0\|_{2 \rightarrow 2} \leq (1 - \mu_D) + K_R$. We need $K_R < \mu_D$.

For PCT at initialisation, $\mu_D \approx e^{-1}/(NM)$ and $K_R \approx (K_f W_m a x^2 R + \rho R)/r_m \text{in}^2$. The condition $K_R < \mu_D$ requires the substrate's bounded-input regime to keep $r_m \text{in}^2 \geq \text{const} \cdot NM$, achievable via RMSNorm.

Open piece: showing $K_R < \mu_D$ *uniformly across L layers* requires the bounded-input regime to be preserved across the cascade – a fixed-point argument that the iterated Lipschitz dynamics stay in the bounded set where (S3) and the $K_R < \mu_D$ condition both hold.

This closes the proof modulo the fixed-point step. **Conjecture 5 is upgraded to a theorem under explicit quantitative conditions on the substrate and the attention diffuseness (S3) preserved across training.** Open: empirical verification of these conditions for PCT trained checkpoints plus tightening the fixed-point argument.